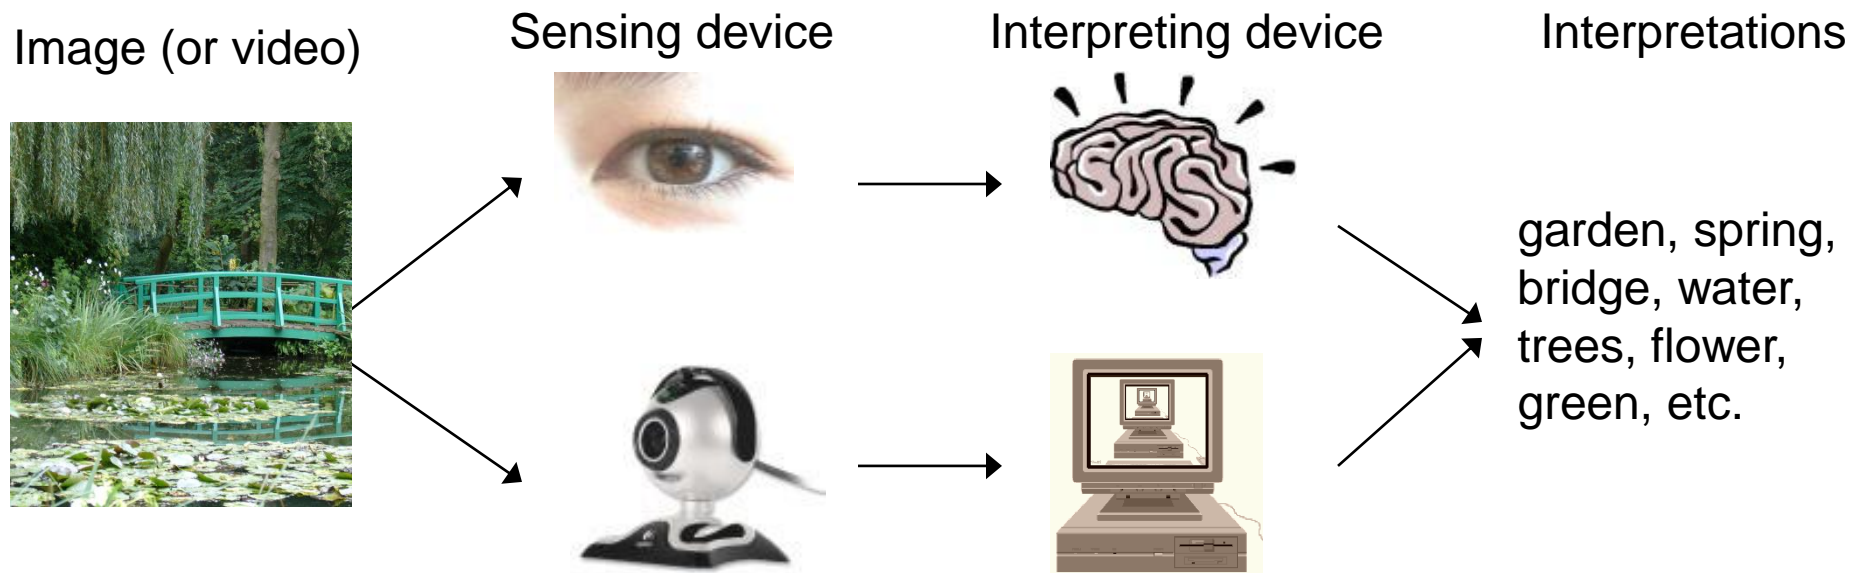


Computer Vision for Robotics

What are “Autonomous Robots”?

- Mobile machines with power, sensing, and computing on-board
- Environments
 - Land (on and under)
 - Water (ditto)
 - Air
 - Space
 - ???

What is (computer) vision?



Computer vision vs human vision



What we see

0	3	2	5	4	7	6	9	8
3	0	1	2	3	4	5	6	7
2	1	0	3	2	5	4	7	6
5	2	3	0	1	2	3	4	5
4	3	2	1	0	3	2	5	4
7	4	5	2	3	0	1	2	3
6	5	4	3	2	1	0	3	2
9	6	7	4	5	2	3	0	1
8	7	6	5	4	3	2	1	0

What a computer sees

What Skills Do Robots Need?

- **Identification:** *What/who is that?*
 - Object detection, recognition
- **Movement:** *How do I move safely?*
 - Obstacle avoidance, homing
- **Manipulation:** *How do I change that?*
 - Interacting with objects/environment
- **Navigation:** *Where am I?*
 - Mapping, localization

Why Vision?

- Pluses


- Rich stream of complex information about the environment
- Primary human sense
- Good cameras are fairly cheap
- Passive → stealthy

- Minuses

- Line of sight only
- Passive → Dependent on ambient illumination

The Vision Problem

How to infer salient
properties of 3-D world
from time-varying 2-D
image projection



Computer Vision Outline

- Image formation
- Image processing
- Motion & Estimation
- Classification & Clustering

Outline: Image Formation

- 3-D geometry
- Physics of light
- Camera properties
 - Focal length
 - Distortion
- Sampling issues
 - Spatial
 - Temporal

Outline: Image Processing

- Filtering
 - Edge
 - Color
 - Shape
 - Texture
- Feature detection
- Pattern comparison

Outline: Motion & Estimation

- Computing temporal image change
 - Magnitude
 - Direction
- Fitting parameters to data
 - Static
 - Dynamic (e.g., tracking)
- Applications
 - Motion Compensation
 - Structure from Motion

Outline: Classification & Clustering

- Categorization
 - Assignment to known groups
- Clustering
 - Inference of group existence from data
 - Special case: Segmentation

Visual Skills: Identification

- Recognizing face/body/structure: *Who/what do I see?*
 - Use shape, color, pattern, other static attributes to distinguish from background, other hypotheses
- Gesture/activity: *What is it doing?*
 - From low-level motion detection & tracking to categorizing high-level temporal patterns
- Feedback between static and dynamic


Minerva Face Detection



Finding people to interact with

Visual Skills: Movement

- Steering, foot placement or landing spot for entire vehicle



MAKRO sewer
shape pattern




Demeter region
boundary detection

UBM Lane & vehicle tracking (with radar)



Visual Skills: Manipulation

- Moving other things
 - Grasping: Door opener (KTH)
 - Pushing, digging, cranes




KTH robot &
typical handle




Clodbusters push a box
cooperatively

Visual Skills: Navigation

- Building a map
- Localization/place recognition
 - *Where are you in the map?*



Laser-based wall map (CMU)




Minerva's ceiling map

Outline: Image Formation

- Geometry
 - Coordinate systems, transformations
 - Perspective projection
 - Lenses
- Radiometry
 - Light emission, interaction with surfaces
- Analog → Digital
 - Spatial sampling
 - Dynamic range
 - Temporal integration

The Image Formation Pipeline



Coordinate System Conventions

- $\mathbf{i}, \mathbf{j}, \mathbf{k}$ unit vectors along positive $\mathbf{k} = \mathbf{i} \times \mathbf{j}$ axes, X, Y, Z respectively;
- Right- vs. left-handed coordinates
- Local coordinate systems: camera, world, etc.




Image Processing Outline

- Images
- Binary operators
- Filtering
 - Smoothing
 - Edge, corner detection
- Modeling, matching
- Scale space

Images

- An image is a matrix of pixels $\mathbf{I}(x, y)$
- Resolution
 - Digital cameras: 1600 X 1200 at a minimum
 - Video cameras: ~640 X 480
- Grayscale: generally 8 bits per pixel \rightarrow Intensities in range [0...255]
- RGB color: 3 8-bit color planes $\mathbf{I}_R, \mathbf{I}_G, \mathbf{I}_B$

Image Conversion

- RGB → Grayscale: Mean color value, or weight by perceptual importance



- Grayscale → Binary: Choose threshold based on histogram of image intensities

Color Representation

- RGB, HSV (hue, saturation, value), YUV, etc.
- Luminance: Perceived intensity
- Chrominance: Perceived color
 - HS(V), (Y)UV, etc.
 - Normalized RGB removes some illumination dependence:

$$r = \frac{R}{R + G + B}, g = \frac{G}{R + G + B}$$

Binary Operations

- Dilation, erosion
 - Dilation: All 0's next to a 1 \rightarrow 1 (Enlarge foreground)
 - Erosion: All 1's next to a 0 \rightarrow 0 (Enlarge background)
- Connected components
 - Uniquely label each n -connected region in binary image
 - 4- and 8-connectedness
- Moments: Region statistics
 - Zeroth-order: Size
 - First-order: Position (centroid)
 - Second-order: Orientation





Image Transformations


- Geometric: Compute new pixel locations
 - Rotate
 - Scale
 - Undistort (e.g., radial distortion from lens)
- Photometric: How to compute new pixel values when $T^{-1}(x', y')$ non-integral
 - Nearest neighbor: Value of closest pixel
 - Bilinear interpolation (2 x 2 neighborhood)
 - Bicubic interpolation (4 x 4)

$$T(x, y) \rightarrow (x', y')$$

Image Comparison: SSD

- Given a template image \mathbf{I}_T and an image \mathbf{I} , how to quantify the similarity between them for a given alignment?
- Sum of squared differences (SSD)

$$\sum_{x,y} [\mathbf{I}_T(x, y) - \mathbf{I}(x, y)]^2$$



Cross-Correlation for Template Matching

- Note that SSD formula can be written:

$$\sum_{x,y} \mathbf{I}_T^2(x, y) + \mathbf{I}^2(x, y) - 2\mathbf{I}_T(x, y)\mathbf{I}(x, y)$$

- First two terms fixed \rightarrow last term measures mismatch—the *cross-correlation*:

$$\sum_{x,y} \mathbf{I}_T(x, y) \cdot \mathbf{I}(x, y)$$

- In practice, normalize by image \mathbf{I} magnitude when shifting template to search for matches

Filtering

- Idea: Analyze neighborhood around some point in image f with filter function h ; put result in new image at corresponding location g
- System properties
 - Shift invariance: Same inputs give same outputs, regardless of location
 - Superposition: Output on sum of images =
 - Sum of outputs on separate images
 - Scaling: Output on scaled image = Scaled output on image
- Linear shift invariance → **Convolution**

Discrete Filtering

- Linear filter: Weighted sum of pixels over rectangular neighborhood—*kernel* defines weights
- Think of kernel as template being matched by correlation
- Convolution: Correlation with kernel rotated 180°
- Dealing with image edges
 - Zero-padding
 - Border replication

1	-1	-1
1	2	-1
1	1	1

Filtering Example 1: $I' = K * I$

K

1	-1	-1
1	2	-1
1	1	1

Rotate

1	1	1
-1	2	1
-1	-1	1


I

2	2	2	3
2	1	3	3
2	2	1	2
1	3	2	2

1	1	1
-1	2	1
-1	-1	1

2	2	2	3
2	1	3	3
2	2	1	2
1	3	2	2

1	1	1		
-1	4	2	2	3
-1	-2	1	3	3
	2	2	1	2
	1	3	2	2



5			


I

I'

1	1	1
-1	2	1
-1	-1	1

2	2	2	3
2	1	3	3
2	2	1	2
1	3	2	2

1	1	1	
-2	4	2	3
-2	-1	3	3
2	2	1	2
1	3	2	2



5	4		


I

I'

1	1	1
-1	2	1
-1	-1	1

	1	1	1
2	-2	4	3
2	-1	-3	3
2	2	1	2
1	3	2	2

I



2	2	2	3
2	1	3	3
2	2	1	2
1	3	2	2


5	4	4	

I'

1	1	1
-1	2	1
-1	-1	1

2	2	2	3
2	1	3	3
2	2	1	2
1	3	2	2

		1	1	1
2	2	-2	6	1
2	1	-3	-3	1
2	2	1	2	
1	3	2	2	



5	4	4	-2


I

I'

1	1	1
-1	2	1
-1	-1	1

2	2	2	3
2	1	3	3
2	2	1	2
1	3	2	2

1	2	2	2	3
-1	4	1	3	3
-1	-2	2	1	2
	1	3	2	2



5	4	4	-2
9			


I

I'

1	1	1
-1	2	1
-1	-1	1

2	2	2	3
2	1	3	3
2	2	1	2
1	3	2	2

2	2	2	3
-2	2	3	3
-2	-2	1	2
1	3	2	2



5	4	4	-2
9	6		

I

I'

Final Result

5	4	4	-2
9	6	14	5
11	7	6	5
9	12	8	5

I'

Smoothing (Low-Pass) Filters

- Replace each pixel with average of neighbors
- Benefits: Suppress noise, aliasing
- Disadvantage: Sharp features blurred
- Types
 - Mean filter (box)
 - Median (nonlinear)
 - Gaussian

$$\frac{1}{9}$$

1	1	1
1	1	1
1	1	1

3 x 3 box filter

Box Filter: Smoothing




Original image



7 x 7 kernel

Gaussian Kernel

- Idea: Weight contributions of neighboring pixels by nearness



$$G_{\sigma} = \frac{1}{2\pi\sigma^2} e^{-\frac{(x^2+y^2)}{2\sigma^2}}$$

Gaussian: Smoothing



Original
image



$\sigma = 1$

7 x 7
kernel



$\sigma = 3$

Gradient

- Think of image intensities as a function $\mathbf{I}(x, y)$
Gradient of image is a vector field as for a normal 2-D height function:

$$\nabla \mathbf{I} = \left(\frac{\partial \mathbf{I}}{\partial x}, \frac{\partial \mathbf{I}}{\partial y} \right)^T = (\mathbf{I}_x, \mathbf{I}_y)^T$$

- **Edge:** Place where gradient magnitude is high;
orthogonal to gradient direction


Edge Causes

- Depth discontinuity
- Surface orientation discontinuity
- Reflectance discontinuity (i.e., change in surface material properties)
- Illumination discontinuity (e.g., shadow)

Edge Detection

- Edge Types


- Step edge (ramp)
- Line edge (roof)




- Searching for Edges:

- **Filter:** Smooth image
- **Enhance:** Apply numerical derivative approximation
- **Detect:** Threshold to find strong edges
- **Localize/analyze:** Reject spurious edges, include weak but justified edges

Sobel Edge Detection: Gradient Approximation



Horizontal



Vertical

Sobel vs. LoG Edge Detection:



Sobel



LoG

Canny Edge Detection

- Derivative of Gaussian
- Non-maximum suppression
 - Thin multi-pixel wide “ridges” down to single pixel
- Thresholding
 - Low, high edge-strength thresholds
 - Accept all edges over low threshold that are connected to edge over high threshold


Canny Edge Detection: Example



(Matlab automatically set thresholds)

Corner Detection

- Basic idea: Find points where two edges meet—i.e., high gradient in orthogonal directions
- Harris corners (Harris & Stephens, 1988), Susan corners (Smith & Brady, 1997)



SUSAN corners

Outline

- Classification terminology
- Unsupervised learning (clustering)
- Supervised learning
 - k -Nearest neighbors
 - Linear discriminants
 - Perceptron, Relaxation, modern variants
 - Nonlinear discriminants
 - Neural networks, etc.
- Applications to computer vision
- Miscellaneous techniques

Classification Terms

- **Data:** A set of N vectors \mathbf{x}
 - Features are parameters of \mathbf{x} ; \mathbf{x} lives in *feature space*
 - May be whole, raw images; parts of images; filtered images; statistics of images; or something else entirely
- **Labels:** C categories; each \mathbf{x} belongs to some c_i
- **Classifier:** Create formula(s) or rule(s) that will assign unlabeled data to correct category
 - Equivalent definition is to parametrize a *decision surface* in feature space separating category members

Key Classification Problems

- What features to use? How do we extract them from the image?
- Do we even have labels (i.e., examples from each category)?
- What do we know about the structure of the categories in feature space?

Unsupervised Learning

- May know number of categories C , but not labels
- If we don't know C , how to estimate?
 - Occam's razor (formalized as Minimum Description Length, or MDL, principle): Favor simpler classifiers over more complex ones
 - Akaike Information Criterion (AIC)
- Clustering methods
 - k-means
 - Hierarchical
 - Etc.

k -means Clustering

- Initialization: Given k categories, N points. Pick k points randomly; these are initial means μ_1, \dots, μ_k
- (1) Classify N points according to nearest μ_i
- (2) Recompute mean μ_i of each cluster from member points
- (3) If any means have changed, goto (1)

Supervised Learning: Assessing Classifier Performance


- Bias: Accuracy or quality of classification
- Variance: Precision or specificity—how stable is decision boundary for different data sets?
 - Related to generality of classification result → Overfitting to data at hand will often result in a very different boundary for new data

Supervised Learning: Procedures

- Validation: Split data into training and test set
 - Training set: Labeled data points used to guide parametrization of classifier
 - % misclassified guides learning
 - Test set: Labeled data points left out of training procedure
 - % misclassified taken to be overall classifier error
- m -fold Cross-validation
 - Randomly split data into m equal-sized subsets
 - Train m times on $m - 1$ subsets, test on left-out subset
 - Error is mean test error over left-out subsets
- Jackknife: Cross-validation with 1 data point left out
 - Very accurate; variance allows confidence measuring

k -Nearest Neighbor Classification

- For a new point, grow sphere in feature space until k labeled points are enclosed
- Labels of points in sphere vote to classify
- Low bias, high variance: No structure assumed



Linear Discriminants

- Basic: $g(\mathbf{x}) = \mathbf{w}^T \mathbf{x} + w_0$
 - \mathbf{w} is weight vector, \mathbf{x} is data, w_0 is bias or threshold weight
 - Number of categories
 - **Two**: Decide c_1 if $g(\mathbf{x}) < 0$, c_2 if $g(\mathbf{x}) > 0$. $g(\mathbf{x}) = 0$ is decision surface—a hyperplane when $g(\mathbf{x})$ linear
 - **Multiple**: Define C functions $g_i(\mathbf{x}) = \mathbf{w}^T \mathbf{x}_i + w_{i0}$. Decide c_i if $g_i(\mathbf{x}) > g_j(\mathbf{x})$ for all $j \neq i$
- Generalized: $g(\mathbf{x}) = \mathbf{a}^T \mathbf{y}$
 - Augmented form: $\mathbf{y} = (1, \mathbf{x}^T)^T$, $\mathbf{a} = (w_0, \mathbf{w}^T)^T$
 - Functions $\mathbf{y}_i = y_i(\mathbf{x})$ can be nonlinear—e.g., $\mathbf{y} = (1, x, x^2)^T$

Dimensionality Reduction

- Functions $\mathbf{y}_i = y_i(\mathbf{x})$ can reduce dimensionality of feature space \rightarrow More efficient classification
- If chosen intelligently, we won't lose much information and classification is easier
- Common methods
 - Principal components analysis (PCA): Maximize total "scatter" of data
$$S_T = \sum_{k=1}^N (x_k - \mu)(x_k - \mu)^T$$
 - Fisher's Linear Discriminant (FLD): Maximize ratio of between-class scatter to within-class scatter

Principal Component Analysis

- Orthogonalize feature vectors so that they are uncorrelated
- Inverse of this transformation takes zero mean, unit variance Gaussian to one describing covariance of data points
- Distance in transformed space is *Mahalanobis distance*
- By dropping eigenvectors of covariance matrix with low eigenvalues, we are essentially throwing away least important dimensions

PCA




FIGURE 2.8. The action of a linear transformation on the feature space will convert an arbitrary normal distribution into another normal distribution. One transformation, \mathbf{A} , takes the source distribution into distribution $N(\mathbf{A}^T \mu, \mathbf{A}^T \Sigma \mathbf{A})$. Another linear transformation—a projection \mathbf{P} onto a line defined by vector \mathbf{a} —leads to $N(\mu, \sigma^2)$ measured along that line. While the transforms yield distributions in a different space, we show them superimposed on the original $x_1 x_2$ -space. A whitening transform, \mathbf{A}_w , leads to a circularly symmetric Gaussian, here shown displaced. From: Richard O. Duda, Peter E. Hart, and David G. Stork, *Pattern Classification*. Copyright © 2001 by John Wiley & Sons, Inc.

Face Recognition

(Belhumeur et al., 1996)

- Given cropped images $\{I\}$ of faces with different lighting, expressions




- Nearest neighbor approach equivalent to correlation (I 's normalized to 0 mean, variance 1)
 - Lots of computation, storage
- PCA projection ("Eigenfaces")
 - Better, but sensitive to variation in lighting conditions
- FLD projection ("Fisherfaces")
 - Best (for this problem)

Local Descriptors

- We know how to detect points
- Next question:

How to *describe* them for matching?




Point descriptor should be:

- 1. Invariant**
- 2. Distinctive**

Invariant Local Features

- Image content is transformed into local feature coordinates that are invariant to translation, rotation, scale, and other imaging parameters

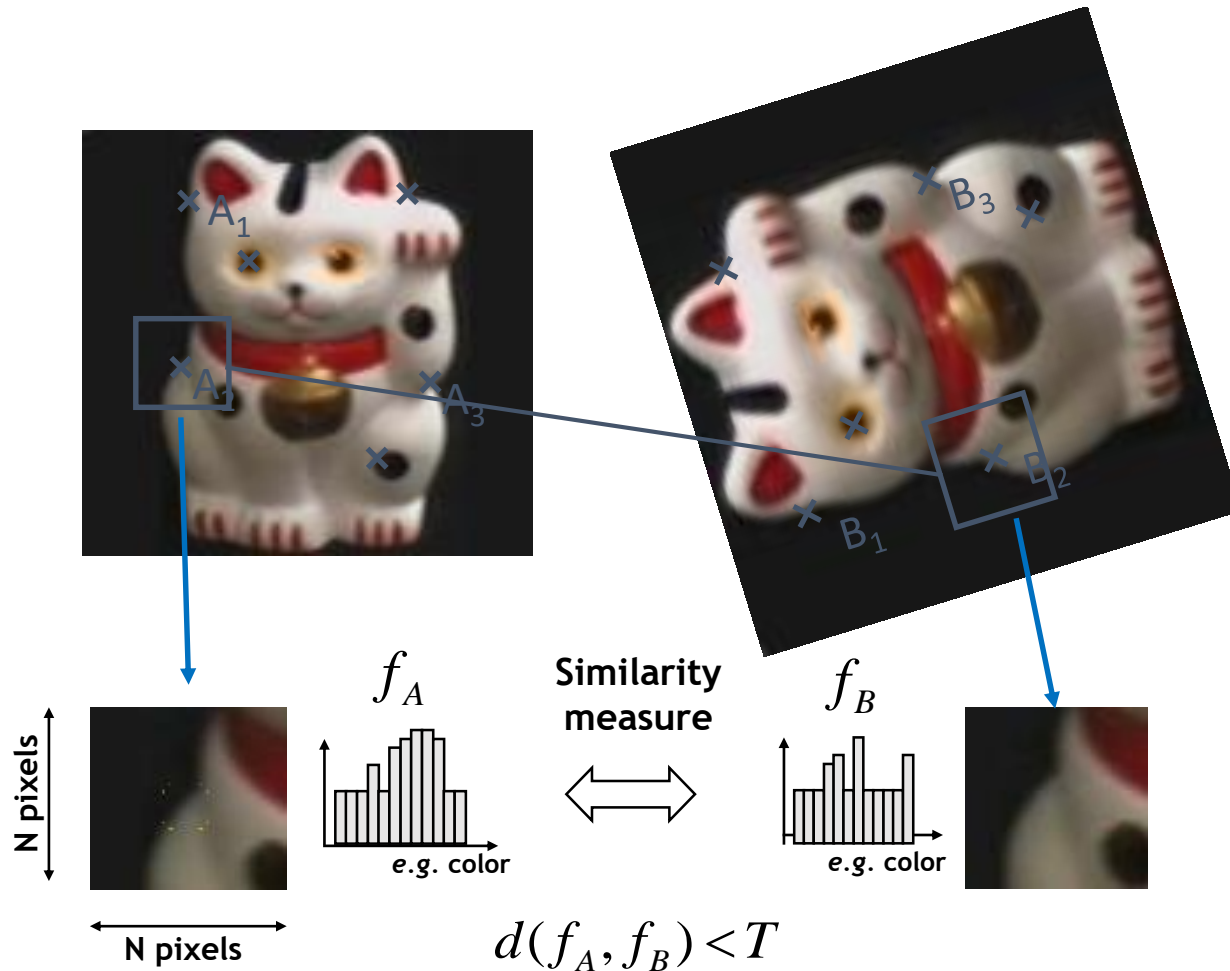


Following slides credit: CVPR 2003 Tutorial on **Recognition and Matching Based on Local Invariant Features** David Lowe

Advantages of invariant local features


- **Locality:** features are local, so robust to occlusion and clutter (no prior segmentation)
- **Distinctiveness:** individual features can be matched to a large database of objects
- **Quantity:** many features can be generated for even small objects
- **Efficiency:** close to real-time performance
- **Extensibility:** can easily be extended to wide range of differing feature types, with each adding robustness

Features Detector: General Approach




1. Find a set of distinctive key-points
2. Define a region around each keypoint
3. Extract and normalize the region content
4. Compute a local descriptor from the normalized region
5. Match local descriptors

SIFT descriptor formation




- Using precise gradient locations is fragile. We'd like to allow some “slop” in the image, and still produce a very similar descriptor
- Create array of orientation histograms (a 4x4 array is shown)
- Put the rotated gradients into their local orientation histograms
 - A gradient's contribution is divided among the nearby histograms based on distance. If it's halfway between two histogram locations, it gives a half contribution to both.
 - Also, scale down gradient contributions for gradients far from the center
- The SIFT authors found that best results were with 8 orientation bins per histogram, **and a 4x4 histogram array.**

Histogram of Oriented Gradients



Visualizing HoG



- a. Average gradient over positive examples
- b. Maximum positive weight in each block
- c. Maximum negative weight in each block
- d. A test image
- e. It's R-HOG descriptor
- f. R-HOG descriptor weighted by positive weights
- g. R-HOG descriptor weighted by negative weights

A simple pipeline - Training

Training Images

